

B561 Assignment 7

Relational Programming

1. Consider the relation schema **Graph**(source INTEGER, target INTEGER) representing a directed graph. The graph *Graph* is *connected* if for each pair of nodes (s, t) in *Graph*, there exists a path in *Graph* from s to t .

An *articulation point* of *Graph* is a node n in *Graph* such that removing the edges in the graph with source n as well as removing the edges with target n results in a graph that is **not** connected. Write a Postgres program that determines the articulation points of *Graph*.

2. Consider the following relational schemas. A tuple $(pid, spid, q)$ is in **Part_SubParts** if *spid* occurs q -times as a **direct** sub-part of *pid*. (For example think of a car that has 4 wheels. Furthermore, then think of a wheel that has 5 bolts.) A tuple (pid, w) is in **Basic_Part** if basic part *pid* has weight w . A basic part is defined as a part that does not have sub-parts. In other words, the pid of the basic part does not occur in the PId column of **Part_SubParts**.

(In the above example, a bolt would be a basic part, but car and wheel would not be basic parts.)

The schemas of **Part_SubParts** and **Basic_Parts** are as follows:

Part_SubParts(PId, SId, Quantity)
Basic_Parts(PId, Weight)

.

You can assume that the domain of each of the attributes in these relations is INTEGER.

Comments:

- (a) In the above example, bolt is a **direct sub-part** of wheel, but not of car. Furthermore, bolt would appear with its weight in **Basic_Parts**, but car nor wheel would appear in this relation.
In other words, only the PId's of parts that have no sub-parts in **Parts_SubParts** are in **Basic_Parts**.
- (b) If the weight of a part is in **Basic_Parts**, the aggregated weight of that part is that weight. Otherwise, the aggregated weight of a part is the sum of the aggregated weights of all its **direct** sub-parts. So the weight function of a part is recursively defined.

Example tables: The following example is based on a desk lamp (pid = 1). Suppose a desk lamp consists of 4 bulbs (pid = 2) and a frame (pid = 3), and a frame consists of a post (pid = 4) and 2 switches (pid = 5).

Furthermore, we will assume that the weight of a bulb is 5, that of a post is 50, and that of a switch is 3.

Then the **Part_SubParts** and **Parts** tables would be as follows:

Parts_SubParts		
PID	SID	Quantity
1	2	4
1	3	1
3	4	1
3	5	2

Parts	
PID	Weight
2	5
4	50
5	3

Then the aggregated weight of a lamp is $4 \times 5 + 1 \times (1 \times 50 + 2 \times 3) = 76$.

Write a Postgres function

Weight(part INTEGER) RETURNS INTEGER AS

..

that takes as input a part-id and returns the aggregated weight for the part with that part-id. In your solution, you can not use cursors nor the **FOR r IN SQL-query** loop statement.

3. Consider the following relational schema. A tuple $(pid, cpid)$ is in **Parent_Child** if pid is a parent of child cid .

```

Parent_Child
-----
|Pid | SId |
-----

```

You can assume that the domain of Pid and SId is INTEGER.

Write a Postgres program that computes the pairs (id_1, id_2) such that id_1 and id_2 belong to the same generation in the **Parent-Child** relation and $id_1 \neq id_2$. (id_1 and id_2 belong to the same generation if their distance to the root in the **Parent-Child** relation is the same.)

4. Suppose you have a weighted directed graph $G = (V, E)$ stored in a ternary table named **Graph** in your database. A triple (n, m, w) in **Graph** indicates that G has an edge (n, m) where n is the source, m is the target, and w is the edge's weight. (In this problem, we will assume that each edge-weight is a positive integer.)

Implement Dijkstra's Algorithm as a Postgres function **Dijkstra** to compute the shortest path lengths (i.e., the distance) from some input node n in G to all other nodes in G . **Dijkstra** should accept an argument n , the source node, and output a table **Paths** which stores the pairs (n, d_m) where d_m is the shortest distance from n to m . To test your procedure, you can use the graph shown in Figure 2. The corresponding table structure for G is given as the following **Graph** table.

Source	Target	Weight
0	1	2
0	4	10
1	2	3
1	4	7
2	3	4
3	4	5
4	2	6

Hint: You can find the details of Dijkstra's Algorithm in the attached pdf document, but you are not required to exactly follow the pseudocode.

When you issue `CALL Dijkstra(0)`, you should obtain the following **Paths** table:

Target	Distance
0	0
1	2
2	5
3	9
4	9